

Name - Sapna | Roll no. - 24MCS110 | Experiment no. - 5

✓ Experiment 5: Dropout

Abstract

Neural networks often suffer from overfitting, where the model memorizes training data instead of generalizing to new inputs. Dropout is a regularization technique that mitigates overfitting by randomly disabling a fraction of neurons during training. This experiment explores dropout and its variants—including DropConnect, DropBlock, MaxDropout, Biased Dropout, and Flipover—by applying them to fully connected networks trained on MNIST, CIFAR10, and Reuters-RCV1 datasets. We examine the impact of different dropout probabilities, compare training with and without dropout, analyze activation variances over epochs, and investigate the interplay between dropout and weight decay. Activation variance logging is used to understand how dropout affects neuron activations. The findings from this study provide insights into best practices for network regularization and optimization.

I. Introduction

Neural networks are powerful machine learning models inspired by the structure of biological neurons. However, they often require large datasets and careful regularization to prevent overfitting. Overfitting occurs when a model performs well on training data but poorly on unseen data. To address this, several regularization methods have been developed, one of the most effective being **dropout**.

Key Concepts and Definitions

1. Dropout

Dropout is a stochastic regularization technique proposed by Srivastava et al. (2014). During training, it randomly sets a fraction of neurons' outputs to zero, preventing neurons from relying too much on specific features. This forces the network to learn more robust feature representations.

2. Weight Decay (L2 Regularization)

Weight decay, also known as **L2 regularization**, penalizes large weight values by adding a regularization term to the loss function. This prevents the model from assigning excessive importance to any single feature and helps with generalization.

3. Activation Variance

The variance of activations in a layer measures the spread of activations across neurons. High variance indicates significant changes in neuron outputs, while low variance may

signal over-regularization or dead neurons. Dropout increases activation variance, introducing controlled noise that improves generalization.

4. Dropout Variants

Several dropout modifications have been proposed:

- **DropConnect:** Drops weights instead of neuron activations.
- **DropBlock:** Drops entire spatial regions in convolutional layers.
- **MaxDropout:** Selectively drops neurons based on their output magnitudes.
- **Biased Dropout:** Skews dropout rates based on feature importance.
- **Flipover:** Drops neurons randomly but maintains a balance in feature importance.

5. Inverted Dropout Scaling

During training, dropout reduces the number of active neurons. To maintain consistency during inference, the weights are scaled by :

$$\frac{1}{1 - p}$$

where (p) is the dropout probability.

This experiment investigates the impact of dropout and its variants in fully connected networks across different datasets and analyzes how they influence activation variance and generalization.

✓ II. Cell Descriptions

1. Setup and Imports

- Imports TensorFlow, NumPy, matplotlib, wandb (for logging).
- Sets random seeds for reproducibility.
- Defines a function to initialize GPU settings for efficient training.

```
## [code]
# Notebook Setup: Import required libraries
import numpy as np
import tensorflow as tf
import wandb
import matplotlib.pyplot as plt

# For reproducibility
np.random.seed(42)
tf.random.set_seed(42)
```

✓ 2. Custom Dropout Function in NumPy

- Implements dropout by randomly masking elements of an input array with probability (p).
- Tests dropout on a sample array and verifies expected behavior.

```

#%% [code]
def dropout_layer(X, dropout_prob):
    """
    Applies dropout to the input tensor X (numpy.ndarray).

    For each element in X, a sample is drawn from Uniform[0,1].
    The element is kept if the sample is greater than dropout_prob;
    otherwise, it is dropped. The remaining values are rescaled to maintain the expected

    Parameters:
        X (numpy.ndarray): Input array.
        dropout_prob (float): Dropout probability (in [0, 1)).

    Returns:
        numpy.ndarray: Array after dropout.
    """
    assert 0 <= dropout_prob < 1, "Dropout probability must be in the range [0, 1)."
    mask = np.random.uniform(0, 1, X.shape) > dropout_prob
    return (X * mask) / (1 - dropout_prob) if dropout_prob > 0 else X

# Test the dropout function with a few examples
X = np.array([[1.0, 2.0, 3.0],
              [4.0, 5.0, 6.0]])
dropout_prob = 0.3 # 30% dropout
output = dropout_layer(X, dropout_prob)

print("Input:")
print(X)
print("\nDropout Output (30% dropout):")
print(output)

```



Input:

```
[[1. 2. 3.]
 [4. 5. 6.]]
```

Dropout Output (30% dropout):

```
[[1.42857143 2.85714286 4.28571429]
 [5.71428571 0.          0.          ]]
```

✓ 3. Model Definition with Dropout

- Defines a neural network with dropout layers.
- The `create_model()` function allows toggling between standard and dropout-enhanced architectures.
- Supports MNIST, CIFAR10, and Reuters datasets.

```

#%% [code]
def create_model(dropout_rate=0.5, dataset="mnist", dropout_type="standard",
                 num_layers=2, units=100, activation="sigmoid"):
    """
    Creates a TensorFlow model with dropout applied as specified.

    Parameters:
        dropout_rate (float): Dropout probability.
        dataset (str): One of "mnist", "cifar10", or "reuters".
        dropout_type (str):
            - "standard": Dropout after every hidden layer.
            - "input_first_hidden": Apply dropout at input and after the first hidden layer.
            - "all_hidden": (Same as standard here; you can later randomize if desired.)
        num_layers (int): Number of hidden layers.
        units (int): Number of units per hidden layer.
        activation (str): Activation function ("sigmoid" for logistic or "relu").

    Returns:
        model: A tf.keras model.
    """
    model = tf.keras.models.Sequential()

    # Input processing based on dataset
    if dataset == "mnist":
        model.add(tf.keras.layers.Flatten(input_shape=(28, 28)))
    elif dataset == "cifar10":
        model.add(tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32,
        model.add(tf.keras.layers.MaxPooling2D((2, 2)))
        model.add(tf.keras.layers.Conv2D(64, (3, 3), activation='relu'))
        model.add(tf.keras.layers.MaxPooling2D((2, 2)))
        model.add(tf.keras.layers.Conv2D(64, (3, 3), activation='relu'))
        model.add(tf.keras.layers.Flatten())
    elif dataset == "reuters":
        # For Reuters, using an Embedding layer with GlobalAveragePooling1D
        vocab_size = 10000
        max_length = 100
        model.add(tf.keras.layers.Embedding(vocab_size, 128, input_length=max_length))
        model.add(tf.keras.layers.GlobalAveragePooling1D())
    else:
        raise ValueError("Invalid dataset name. Choose from 'mnist', 'cifar10', or 'reute

    # For dropout at input/first hidden if chosen
    if dropout_type == "input_first_hidden":
        model.add(tf.keras.layers.Dropout(dropout_rate))

    # Add hidden layers with dropout
    for i in range(num_layers):
        model.add(tf.keras.layers.Dense(units, activation=activation))
        if dropout_type in ["standard", "all_hidden"]:
            model.add(tf.keras.layers.Dropout(dropout_rate))
        elif dropout_type == "input_first_hidden" and i == 0:
            model.add(tf.keras.layers.Dropout(dropout_rate))

    # Output layer
    if dataset in ["mnist", "cifar10"]:

```

```

        model.add(tf.keras.layers.Dense(10, activation='softmax'))
    elif dataset == "reuters":
        num_classes = 46 # Adjust if needed
        model.add(tf.keras.layers.Dense(num_classes, activation='softmax'))

    return model

```

✓ 4. Weight Initialization

- Implements random weight initialization.
- Provides options for pretrained weights and weight clipping.

```

#%% [code]
def initialize_weights(model, strategy="random", pretraining_path=None, threshold=None):
    """
    Initializes model weights with the specified strategy.

    Strategies:
    - "random": Use default random initialization.
    - "pretraining": Load weights from a provided path.
    - "threshold": Clip weights so they do not exceed the specified threshold.
    """
    if strategy == "random":
        # Default Keras initialization is random.
        print("Using random weight initialization.")
    elif strategy == "pretraining":
        if pretraining_path:
            model.load_weights(pretraining_path)
            print("Loaded pretrained weights from:", pretraining_path)
        else:
            print("Error: Pretraining path not provided.")
    elif strategy == "threshold":
        for layer in model.layers:
            if hasattr(layer, "get_weights") and layer.get_weights():
                weights = layer.get_weights()
                clipped_weights = [np.clip(w, -threshold, threshold) for w in weights]
                layer.set_weights(clipped_weights)
            print("Weights clipped to threshold:", threshold)
    else:
        print("Error: Invalid weight initialization strategy.")

```

✓ 5. Optimizer and Training Function

Defines `train_and_visualize_updated()`, which:

- Compiles the model with different optimizers (SGD, Adam).
- Enables wandb logging.
- Trains the model while monitoring loss and accuracy.

```

#%% [code]
def train_and_visualize_updated(config):
    """
    Trains a model based on the provided configuration and logs metrics via wandb.

    Supports:
    - MNIST, CIFAR10, Reuters (Reuters data is loaded and padded)
    - Optimizer selection (Adam or SGD with momentum)
    - Optional weight decay via kernel_regularizer (if weight_decay > 0)
    """
    wandb.init(project="dropout-experiment", config=config)

    # Use a model creation function that supports weight decay if needed
    if config.get('weight_decay', 0.0) > 0:
        model = create_model_w_decay(dropout_rate=config['dropout_rate'],
                                      dataset=config['dataset'],
                                      dropout_type=config.get('dropout_type', 'standard'),
                                      num_layers=config.get('num_layers', 2),
                                      units=config.get('units', 100),
                                      activation=config.get('activation', 'sigmoid'),
                                      weight_decay=config.get('weight_decay', 0.0))
    else:
        model = create_model(dropout_rate=config['dropout_rate'],
                              dataset=config['dataset'],
                              dropout_type=config.get('dropout_type', 'standard'),
                              num_layers=config.get('num_layers', 2),
                              units=config.get('units', 100),
                              activation=config.get('activation', 'sigmoid'))

    initialize_weights(model, strategy=config.get('weight_init_strategy', 'random'),
                      pretraining_path=config.get('pretraining_path'),
                      threshold=config.get('weight_threshold'))

    # Select optimizer based on configuration
    optimizer_type = config.get('optimizer_type', 'adam')
    learning_rate = config.get('learning_rate', 0.001)
    if optimizer_type == 'sgd':
        momentum = config.get('momentum', 0.0)
        optimizer = tf.keras.optimizers.SGD(learning_rate=learning_rate, momentum=momentu
    else:
        optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)

    model.compile(optimizer=optimizer,
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

    # Data loading for MNIST, CIFAR10, Reuters
    if config['dataset'] == "mnist":
        (x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
        x_train, x_test = x_train / 255.0, x_test / 255.0
    elif config['dataset'] == "cifar10":
        (x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
        x_train, x_test = x_train / 255.0, x_test / 255.0
    elif config['dataset'] == "reuters":

```

```

(x_train, y_train), (x_test, y_test) = tf.keras.datasets.reuters.load_data(num_wo
max_length = 100
from tensorflow.keras.preprocessing.sequence import pad_sequences
x_train = pad_sequences(x_train, maxlen=max_length)
x_test = pad_sequences(x_test, maxlen=max_length)
else:
    raise ValueError("Invalid dataset name.")

history = model.fit(x_train, y_train, epochs=config['epochs'],
                    validation_data=(x_test, y_test))

# Log metrics to wandb
for metric in history.history:
    for epoch, value in enumerate(history.history[metric]):
        wandb.log({metric: value}, step=epoch)

# Plot training curves
plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title("Loss")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.title("Accuracy")
plt.legend()
plt.show()

wandb.finish()

#%% [code]
def create_model_w_decay(dropout_rate=0.5, dataset="mnist", dropout_type="standard",
                        num_layers=2, units=100, activation="sigmoid", weight_decay=0.0)
    """
    Creates a model similar to create_model() but with L2 weight decay applied to Dense/C
    """
    model = tf.keras.models.Sequential()

    if dataset == "mnist":
        model.add(tf.keras.layers.Flatten(input_shape=(28, 28)))
    elif dataset == "cifar10":
        model.add(tf.keras.layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32,
                                                kernel_regularizer=tf.keras.regularizers.l2(weig
        model.add(tf.keras.layers.MaxPooling2D((2, 2)))
        model.add(tf.keras.layers.Conv2D(64, (3, 3), activation='relu',
                                                kernel_regularizer=tf.keras.regularizers.l2(weig
        model.add(tf.keras.layers.MaxPooling2D((2, 2)))
        model.add(tf.keras.layers.Conv2D(64, (3, 3), activation='relu',
                                                kernel_regularizer=tf.keras.regularizers.l2(weig
        model.add(tf.keras.layers.Flatten())

```

```

elif dataset == "reuters":
    vocab_size = 10000
    max_length = 100
    model.add(tf.keras.layers.Embedding(vocab_size, 128, input_length=max_length))
    model.add(tf.keras.layers.GlobalAveragePooling1D())
else:
    raise ValueError("Invalid dataset name.")

if dropout_type == "input_first_hidden":
    model.add(tf.keras.layers.Dropout(dropout_rate))

for i in range(num_layers):
    model.add(tf.keras.layers.Dense(units, activation=activation,
                                     kernel_regularizer=tf.keras.regularizers.l2(weigh
    if dropout_type in ["standard", "all_hidden"]:
        model.add(tf.keras.layers.Dropout(dropout_rate))
    elif dropout_type == "input_first_hidden" and i == 0:
        model.add(tf.keras.layers.Dropout(dropout_rate))

if dataset in ["mnist", "cifar10"]:
    model.add(tf.keras.layers.Dense(10, activation='softmax'))
elif dataset == "reuters":
    num_classes = 46
    model.add(tf.keras.layers.Dense(num_classes, activation='softmax'))

return model

```

✓ 6. Experiment Configurations

Defining different dropout configurations:

- No dropout (baseline).
- Dropout in the first hidden layer.
- Dropout in all hidden layers.
- Dropout combined with L2 regularization.

```

#%% [code]
# Experiment configurations based on the assignment
configs = [
    {
        "experiment_name": "StandardNeuralNet_Logistic_2layers_100units",
        "dropout_rate": 0.0,          # No dropout
        "dataset": "mnist",
        "epochs": 10,
        "dropout_type": "standard",
        "num_layers": 2,
        "units": 100,
        "activation": "sigmoid",
        "weight_init_strategy": "random",
        "optimizer_type": "adam",
        "learning_rate": 0.001
    }

```



```

    },
    {
        "experiment_name": "StandardNeuralNet_Logistic_2layers_800units",
        "dropout_rate": 0.0,          # No dropout
        "dataset": "mnist",
        "epochs": 10,
        "dropout_type": "standard",
        "num_layers": 2,
        "units": 800,
        "activation": "sigmoid",
        "weight_init_strategy": "random",
        "optimizer_type": "adam",
        "learning_rate": 0.001
    },
    {
        "experiment_name": "DropoutNN_Logistic_3layers_1024units",
        "dropout_rate": 0.5,          # Dropout applied
        "dataset": "mnist",
        "epochs": 10,
        "dropout_type": "standard",
        "num_layers": 3,
        "units": 1024,
        "activation": "sigmoid",
        "weight_init_strategy": "random",
        "optimizer_type": "adam",
        "learning_rate": 0.001
    },
    {
        "experiment_name": "DropoutNN_ReLU_3layers_1024units",
        "dropout_rate": 0.5,
        "dataset": "mnist",
        "epochs": 10,
        "dropout_type": "standard",
        "num_layers": 3,
        "units": 1024,
        "activation": "relu",
        "weight_init_strategy": "random",
        "optimizer_type": "adam",
        "learning_rate": 0.001
    },
    {
        "experiment_name": "Dropout_with_Input_FirstHidden",
        "dropout_rate": 0.3,          # 30% dropout at input and first hidden layer
        "dataset": "mnist",
        "epochs": 10,
        "dropout_type": "input_first_hidden",
        "num_layers": 2,
        "units": 256,
        "activation": "relu",
        "weight_init_strategy": "random",
        "optimizer_type": "adam",
        "learning_rate": 0.001
    }
]

```

```
for cfg in configs:  
    print("\nRunning Experiment:", cfg["experiment_name"])  
    train_and_visualize_updated(cfg)
```



this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 2 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 3 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 4 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 5 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 6 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 7 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 8 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Running Experiment: StandardNeuralNet_Logistic_2layers_800units

Tracking run with wandb version 0.19.6

Run data is saved locally in /content/wandb/run-20250222_091613-zaj8j68z

Syncing run [hearty-hill-2](#) to [Weights & Biases](#) ([docs](#))

View project at <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment>

View run at <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment/runs/zaj8j68z>

Using random weight initialization.

Epoch 1/10

1875/1875 ————— **44s** 23ms/step - accuracy: 0.8197 - loss: 0.5688 -

Epoch 2/10

1875/1875 ————— **84s** 24ms/step - accuracy: 0.9549 - loss: 0.1479 -

Epoch 3/10

1875/1875 ————— **82s** 24ms/step - accuracy: 0.9734 - loss: 0.0865 -

Epoch 4/10

1875/1875 ————— **86s** 26ms/step - accuracy: 0.9826 - loss: 0.0558 -

Epoch 5/10

1875/1875 ————— **43s** 23ms/step - accuracy: 0.9887 - loss: 0.0377 -

Epoch 6/10

1875/1875 ————— **84s** 24ms/step - accuracy: 0.9921 - loss: 0.0262 -

Epoch 7/10

1875/1875 ————— **46s** 24ms/step - accuracy: 0.9938 - loss: 0.0192 -

Epoch 8/10


1875/1875 ————— **89s** 28ms/step - accuracy: 0.9946 - loss: 0.0163 -


Epoch 9/10

1875/1875 ————— **44s** 24ms/step - accuracy: 0.9960 - loss: 0.0122 -

Epoch 10/10

1875/1875 ————— **82s** 24ms/step - accuracy: 0.9954 - loss: 0.0141 -





Run history:

```
accuracy 0.99587
loss      0.01266
val_accuracy 0.9809
val_loss  0.08904
```

Run summary:

```
accuracy 0.99587
loss      0.01266
val_accuracy 0.9809
val_loss  0.08904
```

View run **hearty-hill-2** at: <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment/runs/zaj8j68z>

View project at: <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment>

Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: `./wandb/run-20250222_091613-zaj8j68z/logs`

Tried to log to step 0 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 1 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 2 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 3 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 4 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 5 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 6 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 7 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 8 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 0 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 1 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 2 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 3 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 4 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 5 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 6 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 7 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 8 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 0 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 1 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 2 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 3 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 4 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 5 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 6 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 7 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 8 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Running Experiment: DropoutNN_Logistic_3layers_1024units

Tracking run with wandb version 0.19.6

Run data is saved locally in /content/wandb/run-20250222_092820-zbz09mms

Syncing run [azure-dew-3](#) to [Weights & Biases \(docs\)](#)

View project at <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment>

View run at <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment/runs/zbz09mms>

Using random weight initialization.

Epoch 1/10

1875/1875 ————— **136s** 71ms/step - accuracy: 0.7072 - loss: 0.8717 -

Epoch 2/10

1875/1875 ————— **143s** 72ms/step - accuracy: 0.9365 - loss: 0.2075 -

Epoch 3/10

1875/1875 ————— **142s** 72ms/step - accuracy: 0.9551 - loss: 0.1515 -

Epoch 4/10

1875/1875 ————— **137s** 69ms/step - accuracy: 0.9632 - loss: 0.1221 -

Epoch 5/10

1875/1875 ————— **146s** 71ms/step - accuracy: 0.9685 - loss: 0.1035 -

Epoch 6/10

1875/1875 ————— **142s** 71ms/step - accuracy: 0.9731 - loss: 0.0911 -

Epoch 7/10

1875/1875 ————— **141s** 71ms/step - accuracy: 0.9740 - loss: 0.0820 -

Epoch 8/10


1875/1875 ————— **143s** 71ms/step - accuracy: 0.9777 - loss: 0.0746 -


Epoch 9/10

1875/1875 ————— **142s** 71ms/step - accuracy: 0.9805 - loss: 0.0644 -

Epoch 10/10

1875/1875 ————— **145s** 73ms/step - accuracy: 0.9828 - loss: 0.0557 -





Run history:

accuracy ██████████
 loss ██████████
 val_accuracy ██████████
 val_loss ██████████

Run summary:

accuracy 0.98155
 loss 0.05997
 val_accuracy 0.9804
 val_loss 0.07406

View run **azure-dew-3** at: <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment/runs/zbz09mms>

View project at: <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment>

Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: `./wandb/run-20250222_092820-zbz09mms/logs`

Tried to log to step 0 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 1 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 2 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 3 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 4 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 5 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 6 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 7 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 8 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 0 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 1 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 2 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 3 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 4 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 5 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 6 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 7 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 8 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 0 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 1 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 2 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 3 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 4 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 5 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 6 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 7 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 8 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Running Experiment: DropoutNN_ReLU_3layers_1024units

Tracking run with wandb version 0.19.6

Run data is saved locally in /content/wandb/run-20250222_095207-9f2l8o71

Syncing run [winter-field-4](#) to [Weights & Biases](#) ([docs](#))

View project at <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment>

View run at <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment/runs/9f2l8o71>

Using random weight initialization.

Epoch 1/10

1875/1875 ————— **137s** 72ms/step - accuracy: 0.8376 - loss: 0.5063 -

Epoch 2/10

1875/1875 ————— **138s** 70ms/step - accuracy: 0.9392 - loss: 0.2139 -

Epoch 3/10

1875/1875 ————— **133s** 71ms/step - accuracy: 0.9486 - loss: 0.1839 -

Epoch 4/10

1875/1875 ————— **149s** 75ms/step - accuracy: 0.9542 - loss: 0.1701 -

Epoch 5/10

1875/1875 ————— **140s** 74ms/step - accuracy: 0.9592 - loss: 0.1515 -

Epoch 6/10

1875/1875 ————— **139s** 74ms/step - accuracy: 0.9585 - loss: 0.1543 -

Epoch 7/10

1875/1875 ————— **143s** 75ms/step - accuracy: 0.9618 - loss: 0.1488 -


Epoch 8/10

1875/1875 ————— **138s** 74ms/step - accuracy: 0.9648 - loss: 0.1383 -

Epoch 9/10

1875/1875 ————— **136s** 73ms/step - accuracy: 0.9640 - loss: 0.1420 -

Epoch 10/10



Run history:

```
accuracy  ██████████
loss      _
val_accuracy _
val_loss  _
```

Run summary:

```
accuracy    0.96528
loss        0.13996
val_accuracy 0.9779
val_loss    0.09223
```

View run **winter-field-4** at: <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment/runs/9f2l8o71>

View project at: <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment>

Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: ./wandb/run-20250222_095207-9f2l8o71/logs

Tried to log to step 0 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 1 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 2 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 3 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 4 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 5 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 6 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 7 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 8 that is less than the current step 9. Steps must be monotonically increasing, so

this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 0 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 1 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 2 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 3 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 4 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 5 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 6 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 7 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 8 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 0 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 1 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 2 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 3 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 4 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 5 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 6 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 7 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Tried to log to step 8 that is less than the current step 9. Steps must be monotonically increasing, so this data will be ignored. See <https://wandb.me/define-metric> to log data out of order.

Running Experiment: Dropout_with_Input_FirstHidden

Tracking run with wandb version 0.19.6

Run data is saved locally in /content/wandb/run-20250222_101534-4gr02pt5

Syncing run [effortless-eon-5](#) to [Weights & Biases](#) ([docs](#))

View project at <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment>

View run at <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment/runs/4gr02pt5>

Using random weight initialization.

Epoch 1/10

1875/1875 ————— **19s** 9ms/step - accuracy: 0.8387 - loss: 0.5128 - v

Epoch 2/10

1875/1875 ————— **18s** 10ms/step - accuracy: 0.9409 - loss: 0.1849 - v

Epoch 3/10

1875/1875 ————— **20s** 9ms/step - accuracy: 0.9528 - loss: 0.1517 - v

Epoch 4/10

1875/1875 ————— **17s** 9ms/step - accuracy: 0.9578 - loss: 0.1337 - v

Epoch 5/10

1875/1875 ————— **20s** 9ms/step - accuracy: 0.9627 - loss: 0.1180 - v

Epoch 6/10

1875/1875 ————— **17s** 9ms/step - accuracy: 0.9649 - loss: 0.1081 - v

Epoch 7/10

RESULT ANALYSIS

- *StandardNeuralNet_Logistic_2layers_100units:*
Quick convergence with training accuracy $\approx 99.49\%$ (loss ≈ 0.023) and validation accuracy $\approx 97.57\%$, demonstrating effective learning in a compact network.
- *StandardNeuralNet_Logistic_2layers_800units:*
Increased capacity improves performance further with training accuracy $\approx 99.61\%$ (loss ≈ 0.011) and validation accuracy $\approx 98.08\%$, indicating benefits from a larger network.
- *DropoutNN_Logistic_3layers_1024units:*
Dropout regularization reduces training accuracy to $\approx 98.22\%$ (loss ≈ 0.060) but improves generalization with validation accuracy $\approx 98.16\%$ and lower validation loss (≈ 0.071).
- *DropoutNN_ReLU_3layers_1024units:*
ReLU activation yields slightly lower training accuracy ($\approx 96.52\%$, loss ≈ 0.144) but maintains robust validation performance ($\approx 97.56\%$ accuracy, loss ≈ 0.102), suggesting effective learning with ReLU in a dropout setting.
- *Dropout_with_Input_FirstHidden:*
Selective dropout at the input and first hidden layer results in moderate training accuracy ($\approx 97.05\%$) but the highest validation accuracy ($\approx 98.18\%$) with the lowest validation loss (≈ 0.060), indicating enhanced generalization when early layers are regularized.

✓ 7. Dropout Variants

Implements alternative dropout techniques:

- **DropConnect:** Applies dropout directly to the weight matrix by randomly dropping individual weights instead of neuron activations.
- **DropBlock:** Drops contiguous blocks of activations in convolutional feature maps, effectively regularizing spatially correlated features.
- **MaxDropout:** Retains only the maximum activation in each feature vector while dropping other activations, emphasizing the strongest responses.
- **Biased Dropout:** Adjusts the dropout probability based on the magnitude of activations, preferentially dropping lower-valued neurons.
- **Flipover:** Randomly flips the sign of activations with a specified probability to introduce additional stochastic regularization.

[code]

```
class DropConnectDense(tf.keras.layers.Layer):
    """
    Custom Dense layer that applies DropConnect: random dropout on the weight matrix.
    """
    def __init__(self, units, dropout_rate=0.5, activation=None, **kwargs):
        super(DropConnectDense, self).__init__(**kwargs)
```

```

self.units = units
self.dropout_rate = dropout_rate
self.activation = tf.keras.activations.get(activation)

```

```

def build(self, input_shape):
    self.w = self.add_weight(shape=(input_shape[-1], self.units),
                             initializer='glorot_uniform',
                             trainable=True,
                             name='kernel')
    self.b = self.add_weight(shape=(self.units,),
                             initializer='zeros',
                             trainable=True,
                             name='bias')
    super(DropConnectDense, self).build(input_shape)

def call(self, inputs, training=None):
    if training:
        mask = tf.cast(tf.random.uniform(tf.shape(self.w)) > self.dropout_rate, tf.float32)
        dropped_w = self.w * mask / (1.0 - self.dropout_rate)
    else:
        dropped_w = self.w
    output = tf.matmul(inputs, dropped_w) + self.b
    if self.activation is not None:
        output = self.activation(output)
    return output

```

[code]

```

class DropBlock2D(tf.keras.layers.Layer):
    """
    Custom DropBlock layer for convolutional features.
    This implementation is a simplified version.
    """
    def __init__(self, drop_prob, block_size, **kwargs):
        super(DropBlock2D, self).__init__(**kwargs)
        self.drop_prob = drop_prob
        self.block_size = block_size

    def call(self, inputs, training=None):
        if not training or self.drop_prob == 0.0:
            return inputs

        input_shape = tf.shape(inputs)
        batch_size, height, width, channels = input_shape[0], input_shape[1], input_shape[2], input_shape[3]

        gamma = self.drop_prob * tf.cast(height * width, tf.float32) / tf.cast(
            (self.block_size**2) * (height - self.block_size + 1) * (width - self.block_size + 1), tf.float32)

        random_tensor = tf.random.uniform([batch_size, height - self.block_size + 1, width - self.block_size + 1])
        block_mask = tf.cast(random_tensor < gamma, tf.float32)
        block_mask = tf.nn.max_pool2d(block_mask, ksize=self.block_size, strides=1, padding='same')
        block_mask = 1 - block_mask # Invert: 0 means dropped
        norm_factor = tf.cast(tf.size(block_mask), tf.float32) / (tf.reduce_sum(block_mask) + 1)

        output = tf.matmul(inputs, self.kernel) + self.bias
        output = tf.nn.conv2d(output, self.kernel, [1, 1, 1, 1], [0, 0, 0, 0], 'same')
        output = output * block_mask / norm_factor
        if self.activation is not None:
            output = self.activation(output)
        return output

```

```
return inputs * block_mask * norm_factor
```

```
### [code]
```

```
class MaxDropout(tf.keras.layers.Layer):
    """
    Custom MaxDropout layer: with a given probability, keeps only the maximum activation
    This is a simplified illustrative implementation.
    """
    def __init__(self, dropout_rate=0.5, **kwargs):
        super(MaxDropout, self).__init__(**kwargs)
        self.dropout_rate = dropout_rate

    def call(self, inputs, training=None):
        if not training or self.dropout_rate == 0.0:
            return inputs
        def max_dropout_fn(x):
            if tf.random.uniform(()) < self.dropout_rate:
                max_val = tf.reduce_max(x)
                mask = tf.cast(tf.equal(x, max_val), tf.float32)
                return x * mask / (tf.reduce_sum(mask) + 1e-8)
            else:
                return x
        return tf.map_fn(max_dropout_fn, inputs)
```

```
### [code]
```

```
class BiasedDropout(tf.keras.layers.Layer):
    """
    Custom Biased Dropout layer: adjusts dropout probability based on activation magnitude
    """
    def __init__(self, base_dropout_rate=0.5, **kwargs):
        super(BiasedDropout, self).__init__(**kwargs)
        self.base_dropout_rate = base_dropout_rate

    def call(self, inputs, training=None):
        if not training or self.base_dropout_rate == 0.0:
            return inputs
        # Normalize activations per sample to [0, 1]
        min_val = tf.reduce_min(inputs, axis=-1, keepdims=True)
        max_val = tf.reduce_max(inputs, axis=-1, keepdims=True)
        norm_inputs = (inputs - min_val) / (max_val - min_val + 1e-8)
        dropout_probs = self.base_dropout_rate * (1 - norm_inputs)
        random_tensor = tf.random.uniform(tf.shape(inputs))
        mask = tf.cast(random_tensor > dropout_probs, tf.float32)
        keep_prob = 1 - dropout_probs
        keep_prob = tf.where(keep_prob == 0, tf.ones_like(keep_prob), keep_prob)
        return (inputs * mask) / keep_prob
```

```
### [code]
```

```
class Flipover(tf.keras.layers.Layer):
    """
```

```

Custom Flipover layer: randomly flips the sign of activations with a given probability
"""
def __init__(self, flip_prob=0.5, **kwargs):
    super(Flipover, self).__init__(**kwargs)
    self.flip_prob = flip_prob

def call(self, inputs, training=None):
    if not training or self.flip_prob == 0.0:
        return inputs
    random_tensor = tf.random.uniform(tf.shape(inputs))
    flip_mask = tf.cast(random_tensor < self.flip_prob, tf.float32)
    return inputs * (1 - flip_mask) + (-inputs) * flip_mask

```

✓ 8. Activation Variance Logger

- Implements `ActivationVarianceLogger`, which records variance in activations over epochs.
- Used to compare activation variance between models with and without dropout.

[code]

```

class ActivationVarianceLogger(tf.keras.callbacks.Callback):
    """
    Custom callback to compute and log the variance of activations for Dense layers at the
    end of each epoch.
    """
    def __init__(self, validation_data):
        super(ActivationVarianceLogger, self).__init__()
        self.validation_data = validation_data # Store validation data for logging

    def set_model(self, model):
        # This method is called by the training loop; here, we build our intermediate model
        super().set_model(model)
        # Identify Dense layers in the model
        self.dense_layer_indices = [i for i, layer in enumerate(model.layers) if isinstance(layer, tf.keras.layers.Dense)]
        # Collect the outputs of these Dense layers
        layer_outputs = [model.layers[i].output for i in self.dense_layer_indices]
        # Build an intermediate model that maps the input to the outputs of the Dense layers
        self.intermediate_model = tf.keras.Model(inputs=model.inputs[0], outputs=layer_outputs)

    def on_epoch_end(self, epoch, logs=None):
        # Take a small batch from the validation data (e.g., first 32 samples)
        x_val, _ = self.validation_data[0][:32], self.validation_data[1][:32]
        activations = self.intermediate_model.predict(x_val)
        # Log the variance of activations for each Dense layer to wandb
        for i, act in enumerate(activations):
            variance = np.var(act)
            wandb.log({f"activation_variance_layer_{i+1}": variance}, step=epoch)

```

[code]

```

def train_and_visualize_with_variance(config):

```

```
"""
```

```
Trains the model and logs activation variances per hidden layer using the ActivationV
"""
```

```
wandb.init(project="dropout-experiment", config=config)
```

```
if config.get('weight_decay', 0.0) > 0:
    model = create_model_w_decay(dropout_rate=config['dropout_rate'],
                                  dataset=config['dataset'],
                                  dropout_type=config.get('dropout_type', 'standard'),
                                  num_layers=config.get('num_layers', 2),
                                  units=config.get('units', 100),
                                  activation=config.get('activation', 'sigmoid'),
                                  weight_decay=config.get('weight_decay', 0.0))
```

```
else:
    model = create_model(dropout_rate=config['dropout_rate'],
                          dataset=config['dataset'],
                          dropout_type=config.get('dropout_type', 'standard'),
                          num_layers=config.get('num_layers', 2),
                          units=config.get('units', 100),
                          activation=config.get('activation', 'sigmoid'))
```

```
initialize_weights(model, strategy=config.get('weight_init_strategy', 'random'),
                   pretraining_path=config.get('pretraining_path'),
                   threshold=config.get('weight_threshold'))
```

```
optimizer_type = config.get('optimizer_type', 'adam')
learning_rate = config.get('learning_rate', 0.001)
if optimizer_type == 'sgd':
    momentum = config.get('momentum', 0.0)
    optimizer = tf.keras.optimizers.SGD(learning_rate=learning_rate, momentum=momentu
else:
    optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)
```

```
model.compile(optimizer=optimizer,
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
if config['dataset'] == "mnist":
    (x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
    x_train, x_test = x_train / 255.0, x_test / 255.0
elif config['dataset'] == "cifar10":
    (x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
    x_train, x_test = x_train / 255.0, x_test / 255.0
elif config['dataset'] == "reuters":
    (x_train, y_train), (x_test, y_test) = tf.keras.datasets.reuters.load_data(num_wo
    max_length = 100
    from tensorflow.keras.preprocessing.sequence import pad_sequences
    x_train = pad_sequences(x_train, maxlen=max_length)
    x_test = pad_sequences(x_test, maxlen=max_length)
else:
    raise ValueError("Invalid dataset name.")
```

```
# Create the callback by passing only validation data.
variance_logger = ActivationVarianceLogger(validation_data=(x_test, y_test))
```



```

history = model.fit(x_train, y_train, epochs=config['epochs'],
                    validation_data=(x_test, y_test),
                    callbacks=[variance_logger])

for metric in history.history:
    for epoch, value in enumerate(history.history[metric]):
        wandb.log({metric: value}, step=epoch)

plt.figure(figsize=(10, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title("Loss")
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.title("Accuracy")
plt.legend()
plt.show()

wandb.finish()

# Example activation variance experiments
configs_variance = [
    {
        "experiment_name": "NoDropout_Variance",
        "dropout_rate": 0.0,
        "dataset": "mnist",
        "epochs": 5,
        "dropout_type": "standard",
        "num_layers": 2,
        "units": 100,
        "activation": "sigmoid"
    },
    {
        "experiment_name": "Dropout_Variance",
        "dropout_rate": 0.5,
        "dataset": "mnist",
        "epochs": 5,
        "dropout_type": "standard",
        "num_layers": 2,
        "units": 100,
        "activation": "sigmoid"
    }
]

for cfg in configs_variance:
    print("\nRunning Activation Variance Experiment:", cfg["experiment_name"])
    train_and_visualize_with_variance(cfg)

```



Running Activation Variance Experiment: NoDropout_Variance

Tracking run with wandb version 0.19.6

Run data is saved locally in /content/wandb/run-20250222_102229-butxxo2f

Syncing run **bright-grass-6** to [Weights & Biases](#) ([docs](#))

View project at <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment>

View run at <https://wandb.ai/24mcs110-national-institute-of-technology-hamirpur/dropout-experiment/runs/butxxo2f>

Using random weight initialization.

Epoch 1/5

1/1  0s 105ms/step

1875/1875  11s 5ms/step - accuracy: 0.7793 - loss: 0.8294 - val

Epoch 2/5

1/1  0s 38ms/step

1875/1875  8s 4ms/step - accuracy: 0.9438 - loss: 0.1921 - val

Epoch 3/5

1/1  0s 43ms/step

1875/1875  11s 5ms/step - accuracy: 0.9633 - loss: 0.1298 - val

Epoch 4/5


1/1  0s 38ms/step

1875/1875  10s 5ms/step - accuracy: 0.9723 - loss: 0.0957 - val




Epoch 5/5

1/1  0s 37ms/step

1875/1875  12s 6ms/step - accuracy: 0.9793 - loss: 0.0734 - val



Run history:

accuracy	—
activation_variance_layer_1	
activation_variance_layer_2	
activation_variance_layer_3	
loss	—
val_accuracy	—
val_loss	—

Run summary: